


The real context-vector patch does not separate from the floor controls

$f_{CV} = +0.023$ = mean over (context, layer, scope) of $[p_{up} - self_{c0}]$; $P \uparrow$ moves to the base-CV null, $P \downarrow$ falls to the random/norm-matched floor — no input localization



base-CV null (00) trained-CV null (0+) random-CV floor
trained CV $\rightarrow$ base (real patch) base CV $\rightarrow$ trained (real patch) norm-matched floor