

The realized write does not point toward the training-data target

Mean $\cos(\text{write}, \delta)$ per behavior; positive-only $\approx$ contrastive everywhere, so the negatives do not rotate the write direction.

