

The context-vector patch carries almost none of the finetuning shift

Cross-model patch at layer 10, answer-side read at layer 14; 16 cells $\times$ 280 pairs each

