


v_C predicts early answer positions best, decays with p



v_C = last-input-token context vector (mean over 48 probes); v_A = mean over all answer tokens. LOCO ridge skill-over-mean R^2 , PCA target dim = $\min(48, n-2)$. n_{contexts} at $p=0$ is 50, at $p=64$ is 50 (L18). Dashed line: $v_C \rightarrow v_A$ reference at L18.