

$v_C \rightarrow v_A$ per-layer predictive skill: linear vs kernel, two v_C recipes

skill over predict-the-mean
(held-out LOCO R^2)

