

Predicting behavior B from v_A vs $M v_C$ (ridge readout)

best-layer held-out Spearman ρ , 95% CI (2000-bootstrap over 50 contexts) · Qwen2.5-7B · n=50 · LOCO

