


Two independent reads agree

corrected on-policy read vs #664's in-loop probe: Spearman $\rho=0.92$, Pearson $r=0.99$, $p<0.001$ ($n=16$)

corrected on-policy read, trained - base (nat)



in-loop teacher-forced band-stop read (nat)

issue #734 corrected_reread + #664 inloop_band_stop