

Sycophancy: finetuning shift predicts trait expression ($r = 0.85$)

$n=24$ finetunes; Pearson $r = 0.85$ [95% CI 0.72, 0.93] at the best of 28 layers

Sycophancy trait-expression score (0-100, graded judge)

