

Finetuning shift: trait-expression score vs projection — persona vector (top) vs one norm-matched random direction (bottom); selected layers evil L23, syco L20, hall L25

