

Context→answer map: held-out R^2 vs training-set size

Qwen-2.5-7B-Instruct, L19, last-token → mean-answer; fixed 1,000-ctx test.

963k point: mixed LMSYS+WildChat; KRR $\geq 150k$ = Nyström $m=16,384$.

Dash-dot: single-draw reliability ceiling (0.947).

