

Context→answer map: held-out R^2 vs training-set size

Qwen-2.5-7B-Instruct, L19, last-token → mean-answer; fixed 1,000-ctx test set.

≥150k: 963k point mixed LMSYS+WildChat; KRR = Nyström $m=16,384$ (validated at 50k). Dash-dot: reliability ceiling 0.947.

