

Context→answer map: held-out R^2 vs training-set size

Qwen-2.5-7B-Instruct, LMSYS; last-token → mean-answer, layer 19; fixed 1,000-ctx test set.
Dashed: n=10,000 refit (extended λ grid, same test set).

