

Context $\rightarrow$ answer map: held-out R^2 vs training-set size

Qwen-2.5-7B-Instruct, LMSYS; last-token $\rightarrow$ mean-answer, layer 19; fixed 1,000-ctx test set.
Dashed: $n=10,000$ refit; $n=50,000$ plan-B refit (squares, bootstrap 95% CI; single fit per predictor).

