

In-sample training reconstruction R^2 vs shuffled-pairing null (R^2 0.83-0.86)

linear map beats the nonlinear map; both far above the shuffled-pairing null (in-sample, not held-out)

