

The prompt state linearly predicts the answer's mean activation profile (R^2 0.83-0.86)

linear map beats the nonlinear map; both far above the shuffled-context null

