



Per-unit data behind the R3 in-sample reconstruction R^2 / cosine (ridge map)

top: per-training-context reconstruction cosine (behind mean cosine); bottom: per-output-dimension R^2 (behind the pooled R^2). Black line = the aggregate.

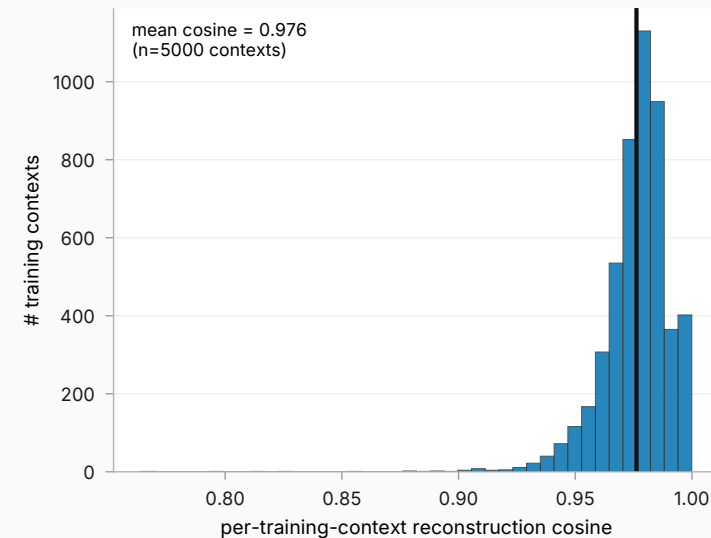
Evil — per-context recon cosine



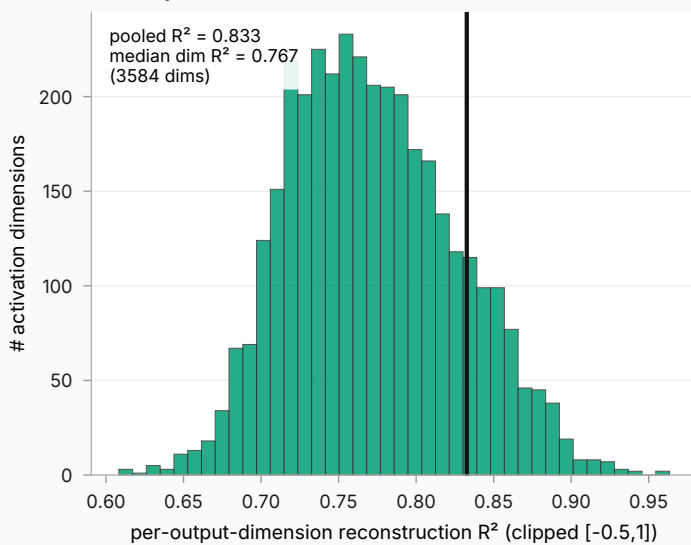
Sycophancy — per-context recon cosine



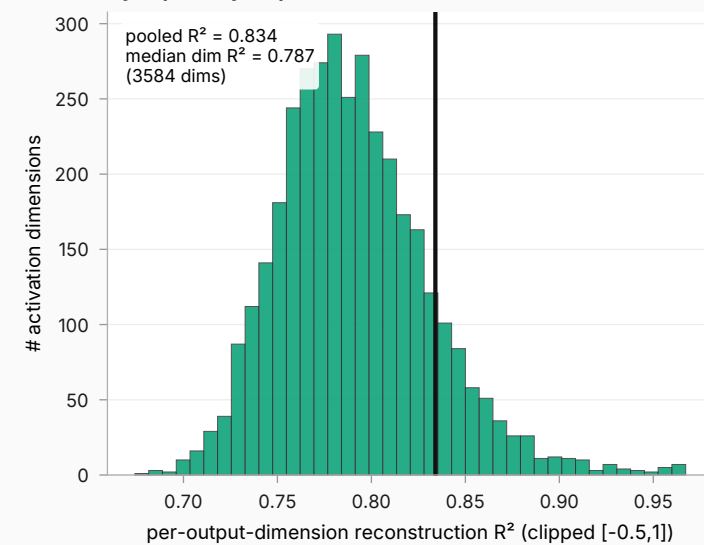
Hallucination — per-context recon cosine



Evil — per-dimension R^2



Sycophancy — per-dimension R^2



Hallucination — per-dimension R^2

