

Per-layer predictivity: sycophancy

LOCO ridge to graded expression, $n = 50$ contexts

