

Base weights over the trained models' own answers reproduce the on-policy chain: all nine paired differences are consistent with zero

