


per-question grain, frozen layer 25 (LOCO, n=1,994 rows / 50 contexts)



where the CoT gain concentrates

