

Predictability vs token-level surprise per position

Each point = one answer position t (1-16, 32, 64, 128)

