

Per-context predictability behind the $t = 128$ recovery read

Own vs mismatched answer arms; test contexts with 144-token spans

