

Predicting later answers from earlier context states

answer-k prediction; 95% bootstrap CIs, n=500 conversations

held-out skill (R^2 vs corpus-mean baseline)

