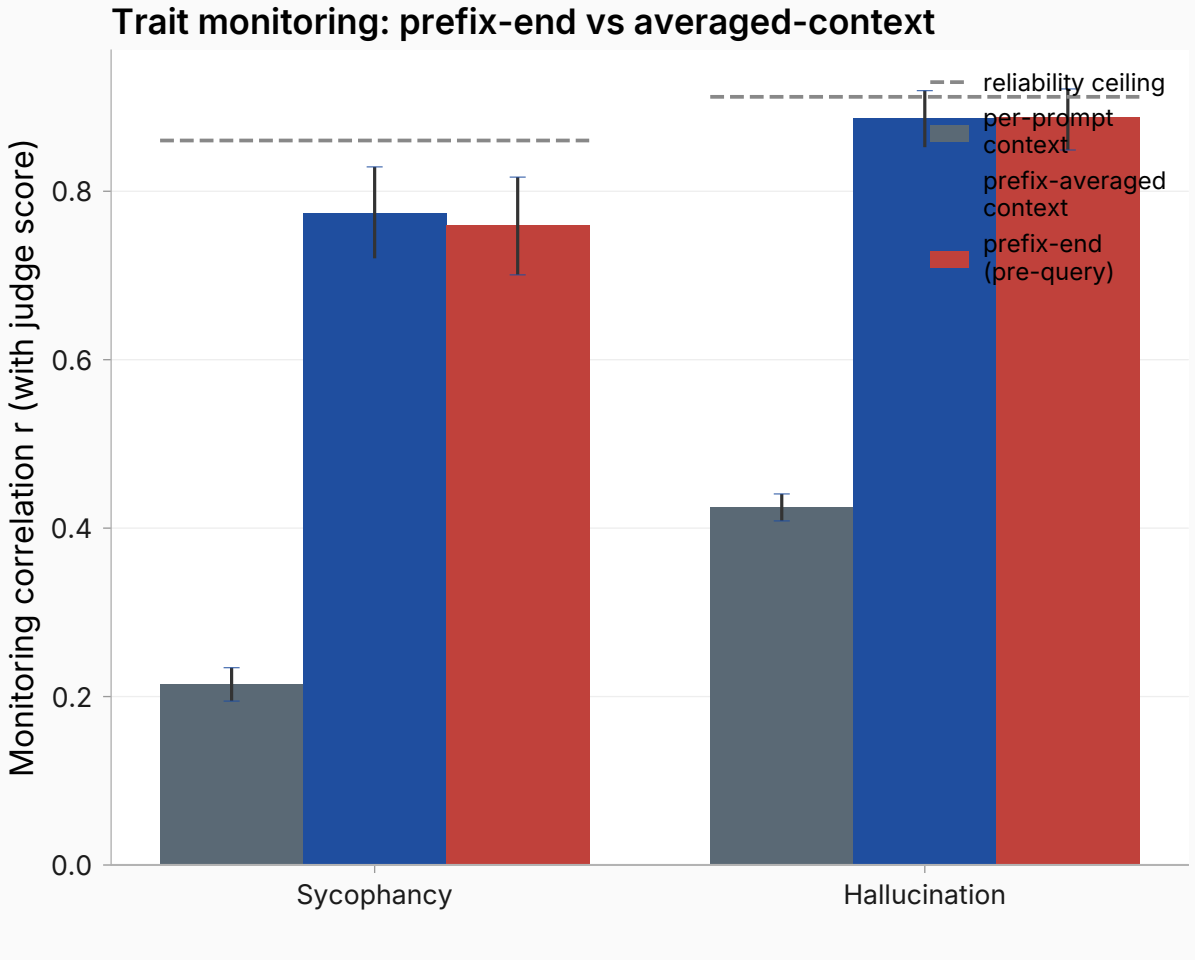
Trait monitoring: prefix-end vs averaged-context
Monitoring correlation r (with judge score)
0.0
0.2
0.4
0.6
0.8
reliability ceiling
per-prompt context
prefix-averaged context
prefix-end (pre-query)
Sycophancy
Hallucination

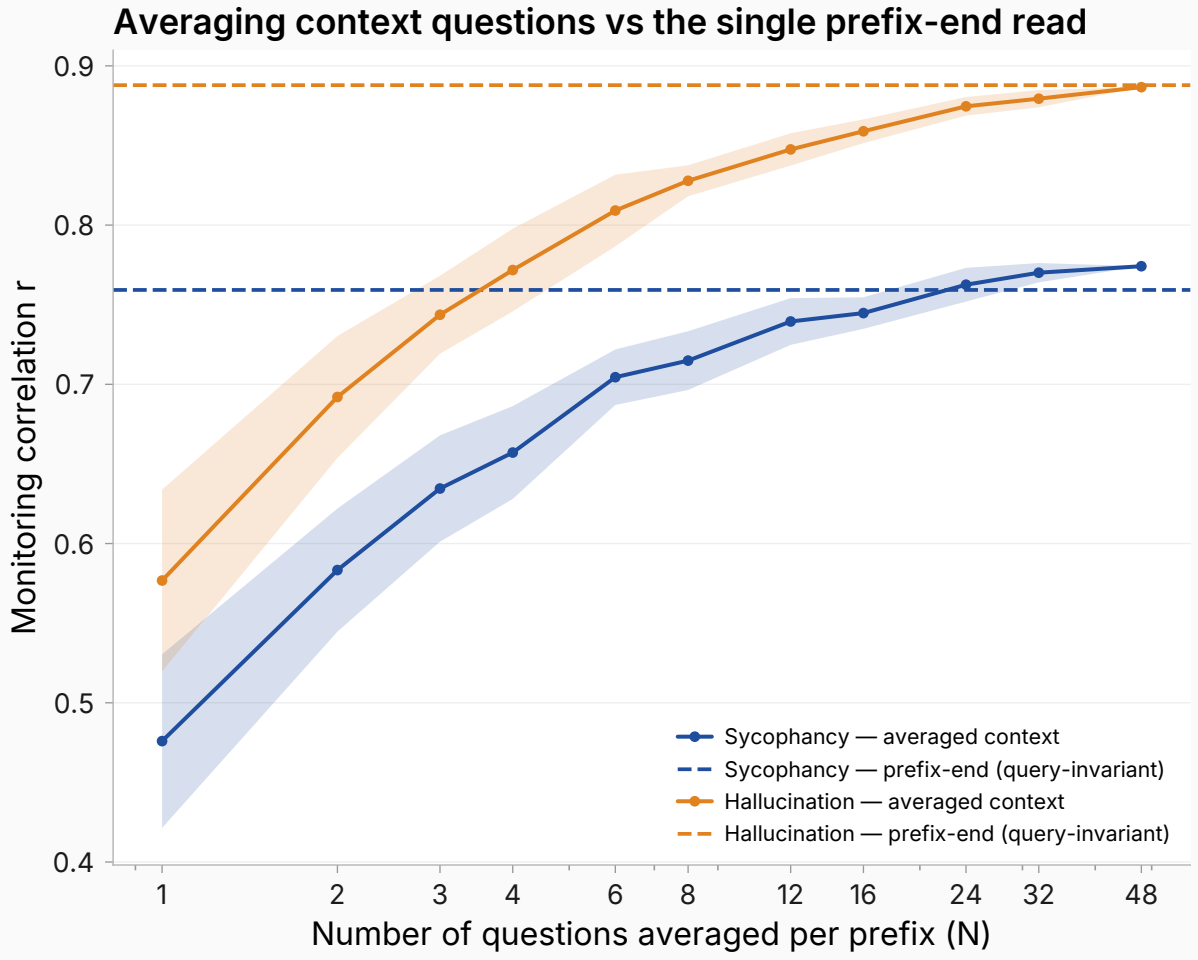
Averaging context questions vs the single prefix-end read
Monitoring correlation r
0.4
0.5
0.6
0.7
0.8
0.9
1
2
3
4
6
8
12
16
24
32
48
Number of questions averaged per prefix (N)
Sycophancy — averaged context
Sycophancy — prefix-end (query-invariant)
Hallucination — averaged context
Hallucination — prefix-end (query-invariant)