

Trait monitoring from one pre-query vector vs question-averaged context reads
(#1092 dense core: 149 real conversation prefixes × 48 queries, instruct model, layer 14)

