





- - $R^2 = 0$ · predict the training mean
 - identity+bias $\hat{y} = x + b$ (lower panel)
 - 0: direct transfer
 - 6: reparam contexts
 - 7: reparam answers
 - 8: reparam both
 - cross: source ctx → target answer (fresh fit)
 - within-model ceiling (target's own map)