



Alignment maps between the two checkpoints — a DIFFERENT regression, not comparable to the panel below



Every forward pair of the Tülu-3 ladder: how much of the context→answer map survives
median over 7 non-degenerate corpora, every corpus overplotted



- - $R^2 = 0$ · predict the training mean
 - - shuffled-pairing null, per tier (t0/t6 drawn; t7/t8 = -0.0008 $\approx$ the zero line)
 - - identity+bias $\hat{y} = x + b$ (lower panel) — MEASURED
 ○ identity+bias — BORROWED from the target's round-3 pairs
 ● 0: direct transfer
 ● 6: reparam contexts
 ● 7: reparam answers
 ● 8: reparam both
 ◆ cross: source ctx → target answer (fresh fit)
 ◆ within-model ceiling (target's own map)