



Alignment maps between the two checkpoints — a DIFFERENT regression, not comparable to the panel below



Every forward pair of the Tülu-3 ladder: how much of the context→answer map survives
median over 7 non-degenerate corpora, every corpus overplotted



-- $R^2 = 0$ · predict the training mean
 identity+bias $\hat{y} = x + b$ (lower panel)
 0: direct transfer
 6: reparam contexts
 7: reparam answers
 8: reparam both
 cross: source ctx → target answer (fresh fit)
 within-model ceiling (target's own map)