



A DIFFERENT regression, plotted apart so it is not read against the panel below: identity+bias  $\hat{y} = x + b$  aligning the two checkpoints' summary vectors (target contexts→source contexts; source answers→target answers). The ceiling below maps context→answer INSIDE one model — a high value here is not "the baseline beats the self map".



Every forward pair of the Tülu-3 ladder: how much of the context→answer map survives  
median over 7 non-degenerate corpora, every corpus overplotted



— shuffled-pairing null, per tier (t0/t6 drawn; t7/t8 = -0.0008 ≈ the zero line)

— identity+bias baseline  $\hat{y} = x + b$  (lower panel)

— 0: direct transfer

— 6: reparam contexts

— 7: reparam answers

— 8: reparam both

— cross: source ctx → target answer (fresh fit)

— within-model ceiling (target's own map)