

Per eval dataset: every forward stage pair, at each reparameterization tier

--- shuffled-pairing null, per tier (t0/t6 lower panel; t7/t8 ≈ the zero line)
- identity+bias baseline $\hat{y} = x + b$ (lower panel)
- cross: source ctx → target answer (fresh fit)
- within-model ceiling
- 0: direct transfer
- 6: reparam contexts
- 7: reparam answers
- 8: reparam both

