

Per eval dataset: every forward stage pair, at each reparameterization tier

- - shuffled-pairing null, per tier (t0/t6 lower panel; t7/t8 $\approx$ the zero line)
 - - identity+bias baseline $\hat{y} = x + b$ (lower panel)
 - - cross: source ctx $\rightarrow$ target answer (fresh fit)
 - - within-model ceiling
 0: direct transfer
 6: reparam contexts
 7: reparam answers
 8: reparam both

