





- $R^2 = 0$ · predict the training mean
- identity+bias $\hat{y} = x + b$ (lower panel)
- 0: direct transfer
- 1: context offset
- 2: answer offset
- 3: bias offset
- 4: global scaling
- 5: rotation + bias
- 6: reparam contexts
- 7: reparam answers
- 8: reparam both
- cross: source ctx → target answer (fresh fit)
- within-model ceiling (target's own map)
- source self-map (source's own ceiling)