





-- $R^2 = 0$ · predict the training mean
 identity+bias $\hat{y} = x + b$ (lower panel)
 0: direct transfer
 5: rotation
 6: reparam contexts
 7: reparam answers
 cross: source ctx → target answer (fresh fit)
 within-model ceiling (target's own map)