





--  $R^2 = 0$  · predict the training mean  
 identity+bias  $\hat{y} = x + b$  (lower panel)  
 0: direct transfer  
 5: rotation  
 6: reparam contexts  
 7: reparam answers  
 cross: source ctx → target answer (fresh fit)  
 within-model ceiling (target's own map)