

Does the AMOUNT of training in between predict how badly the context→answer map transfers?
three ways to count it — the row that matters is whichever one the ladder does not confound

