


Does a training stage disrupt the context→answer map more on its OWN training data?

Within each corpus: direct transfer, one point per stage pair



Naive gap → gap with pair and corpus effects removed

