


Does the training BETWEEN two models disrupt the context→answer map more on the data that training used?

Within each corpus: direct transfer, one point per model pair



Naive gap → gap with pair and corpus effects removed

