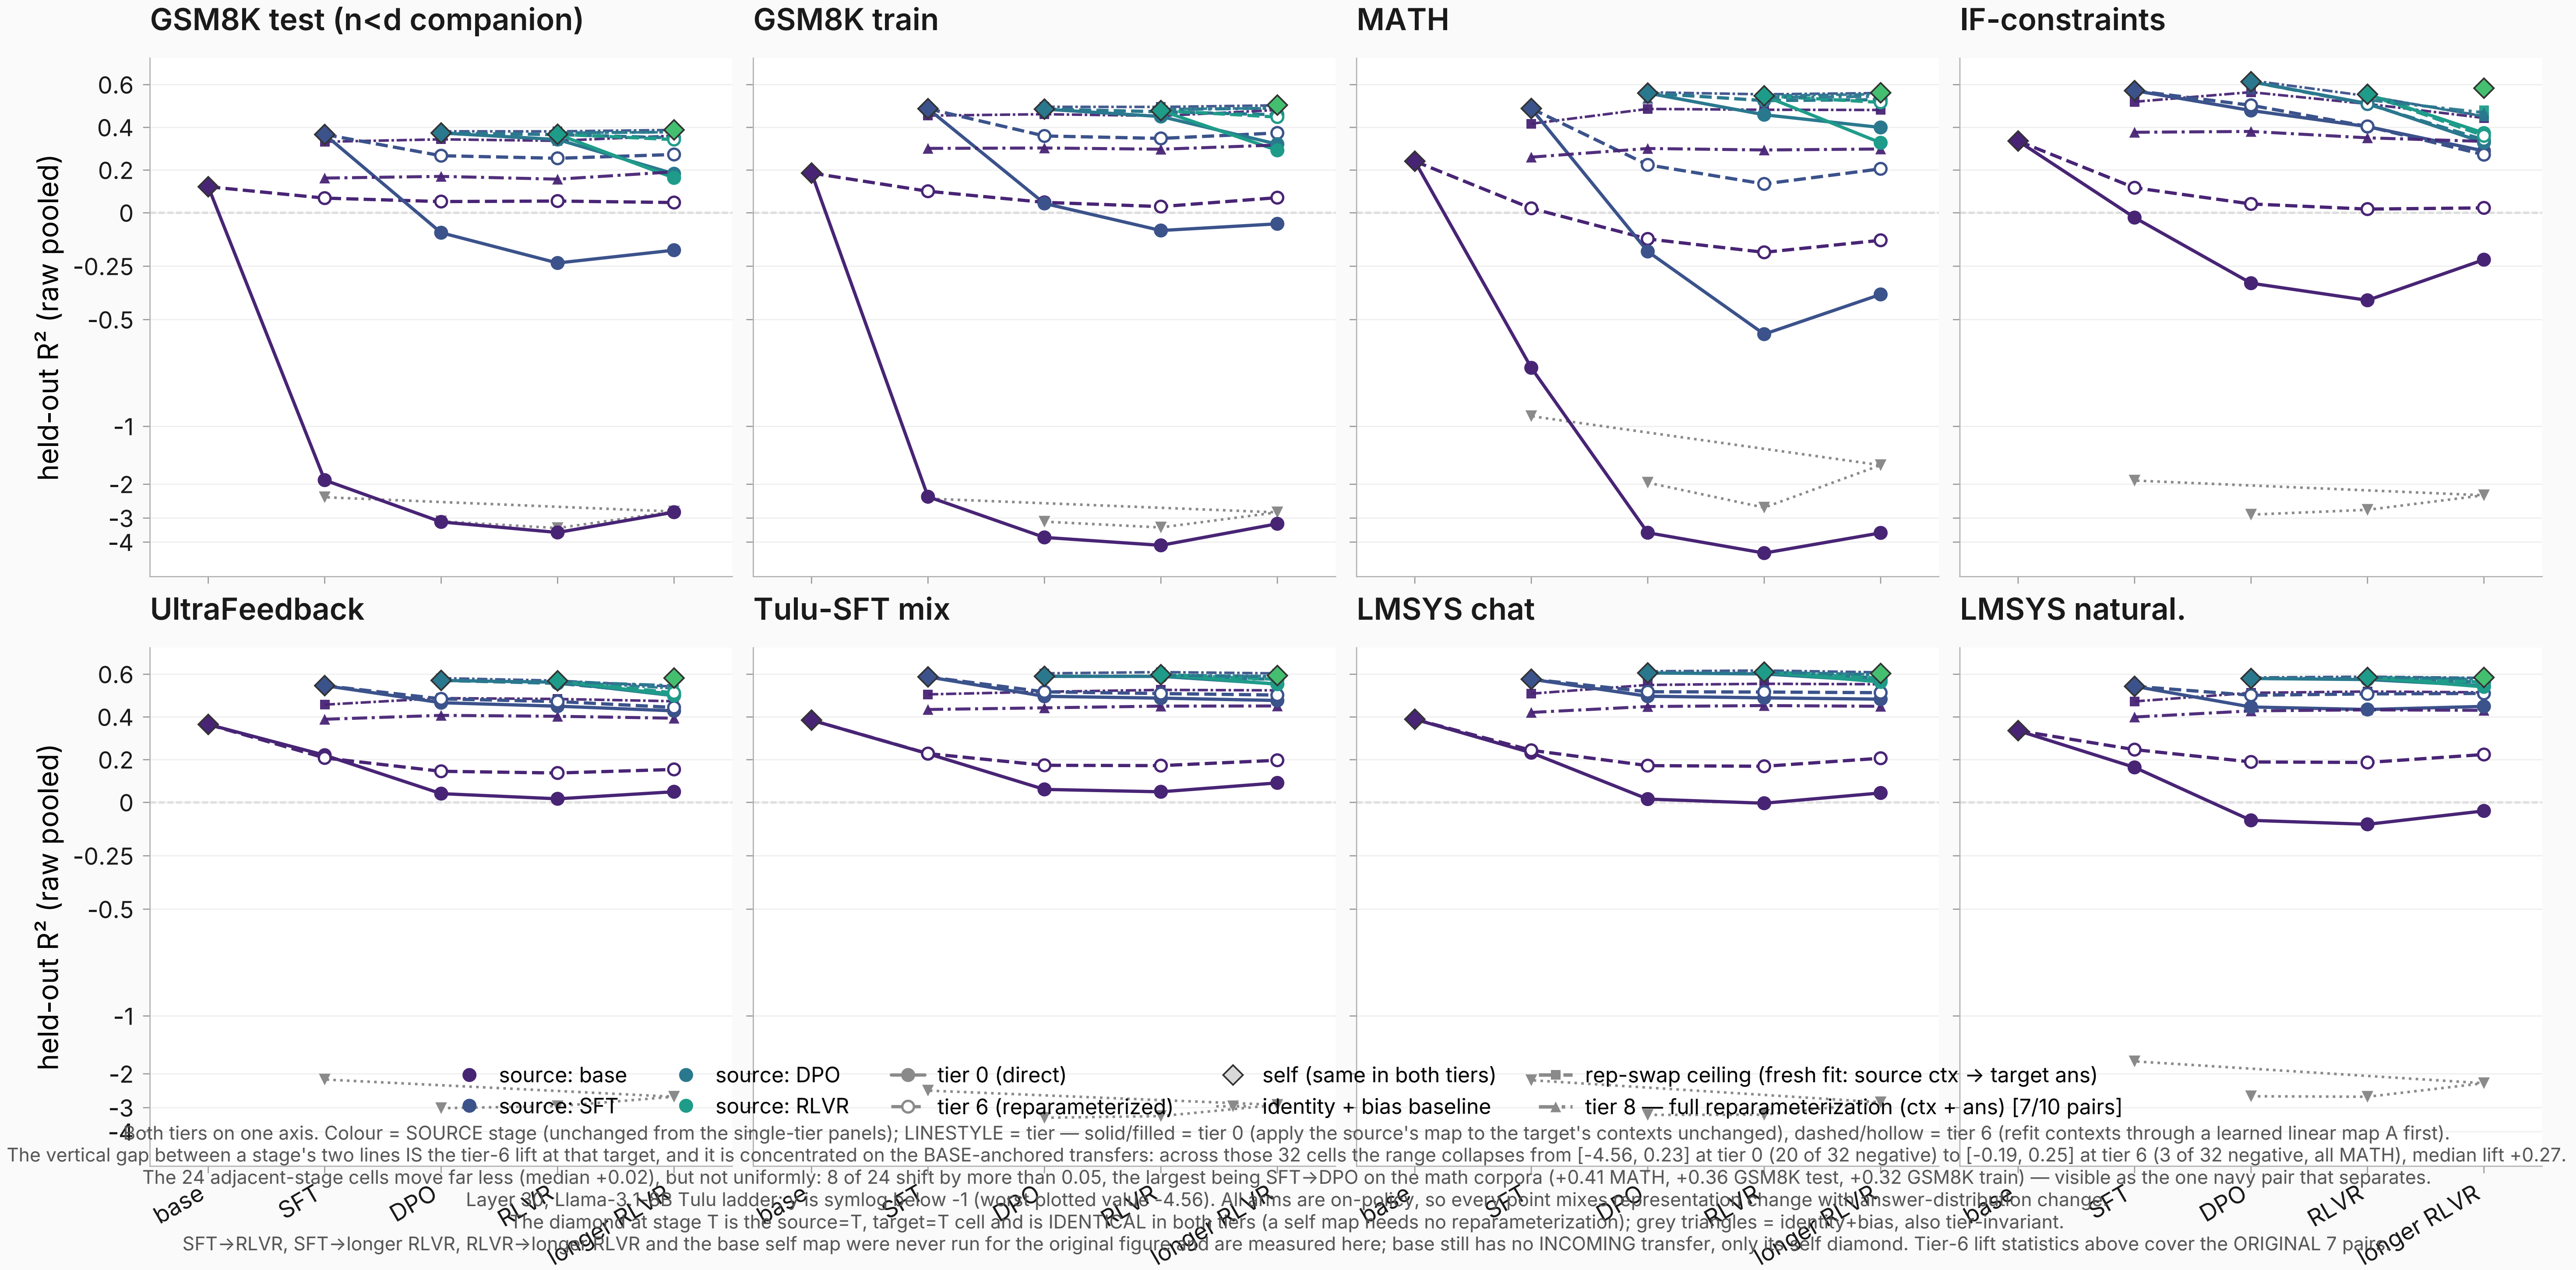

Both tiers on one axis. Colour = SOURCE stage (unchanged from the single-tier panels); LINESTYLE = tier — solid/filled = tier 0 (apply the source's map to the target's contexts unchanged), dashed/hollow = tier 6 (refit contexts through a learned linear map A first).
The vertical gap between a stage's two lines IS the tier-6 lift at that target, and it is concentrated on the BASE-anchored transfers: across those 32 cells the range collapses from [-4.56, 0.23] at tier 0 (20 of 32 negative) to [-0.19, 0.25] at tier 6 (3 of 32 negative, all MATH), median lift +0.27.
The 24 adjacent-stage cells move far less (median +0.02), but not uniformly: 8 of 24 shift by more than 0.05, the largest being SFT→DPO on the math corpora (+0.41 MATH, +0.36 GSM8K test, +0.32 GSM8K train) — visible as the one navy pair that separates.
Layer 30, Llama-3.1-8B Tulu ladder; y is symlog below -1 (worst plotted value -4.56). All arms are on-policy, so every point mixes representation change with answer-distribution change.
The diamond at stage T is the source=T, target=T cell and is IDENTICAL in both tiers (a self map needs no reparameterization); grey triangles = identity+bias, also tier-invariant.
SFT→RLVR, SFT→longer RLVR, RLVR→longer RLVR and the base self map were never run for the original figure and are measured here; base still has no INCOMING transfer, only its self diamond. Tier-6 lift statistics above cover the ORIGINAL 7 pairs.