


Tier 0 — direct transfer (no correction)

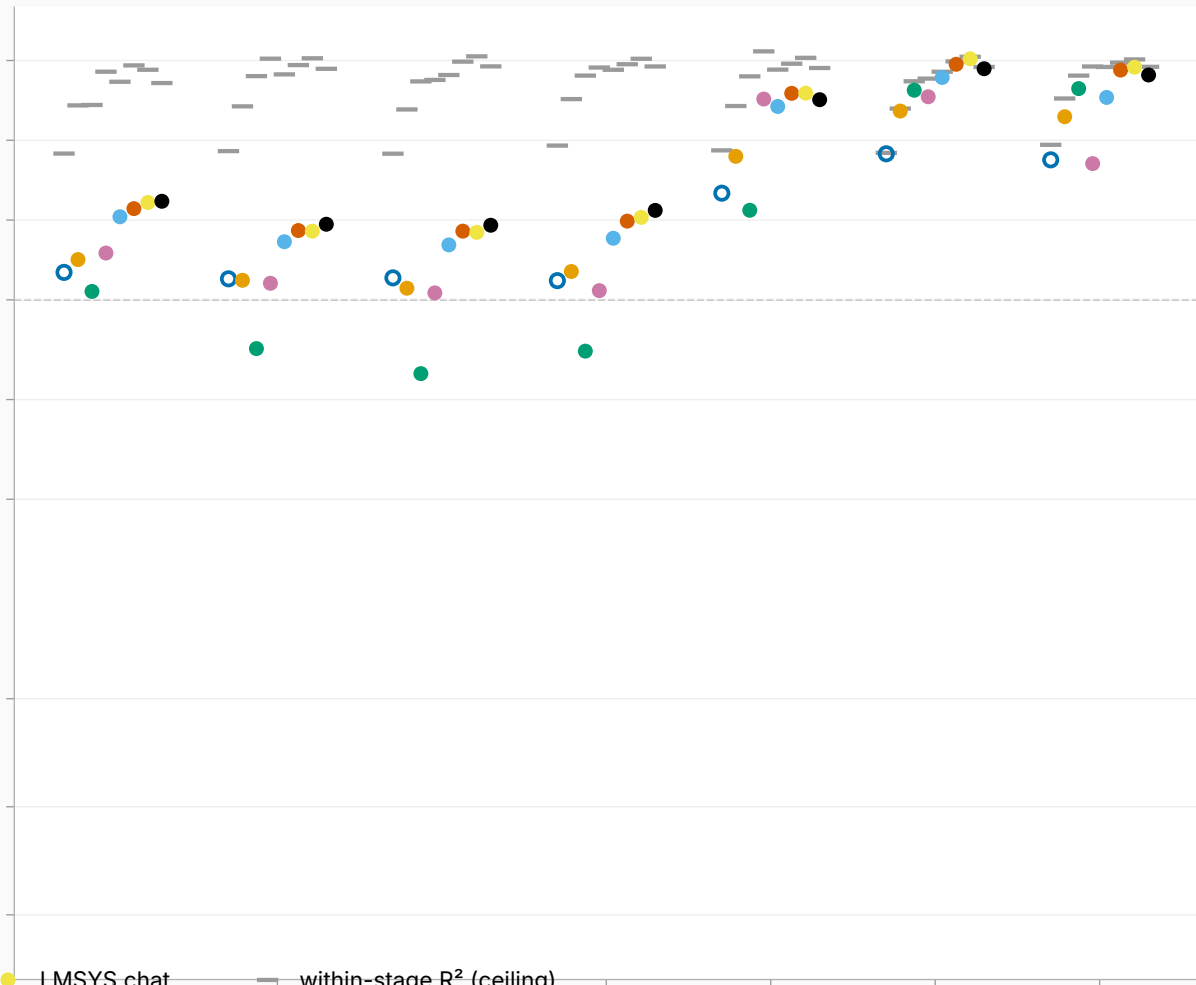
apply stage i's map to stage j's contexts unchanged

held-out R^2 on target pairs (raw pooled)



Tier 6 — linear reparameterization of contexts

refit contexts through a learned linear map A , then apply stage i's map: $\hat{y} = W_i(A \times)$



Layer 30 of llama-3.1-8B Tulu ladder. One point per corpus. One symbol below -1 (worst transfer -4.56) follows marker for GSM8K test, the 4-d companion cell.
All arms are on-policy — each stage answers in its own words — so every point mixes representation change with answer distribution change.