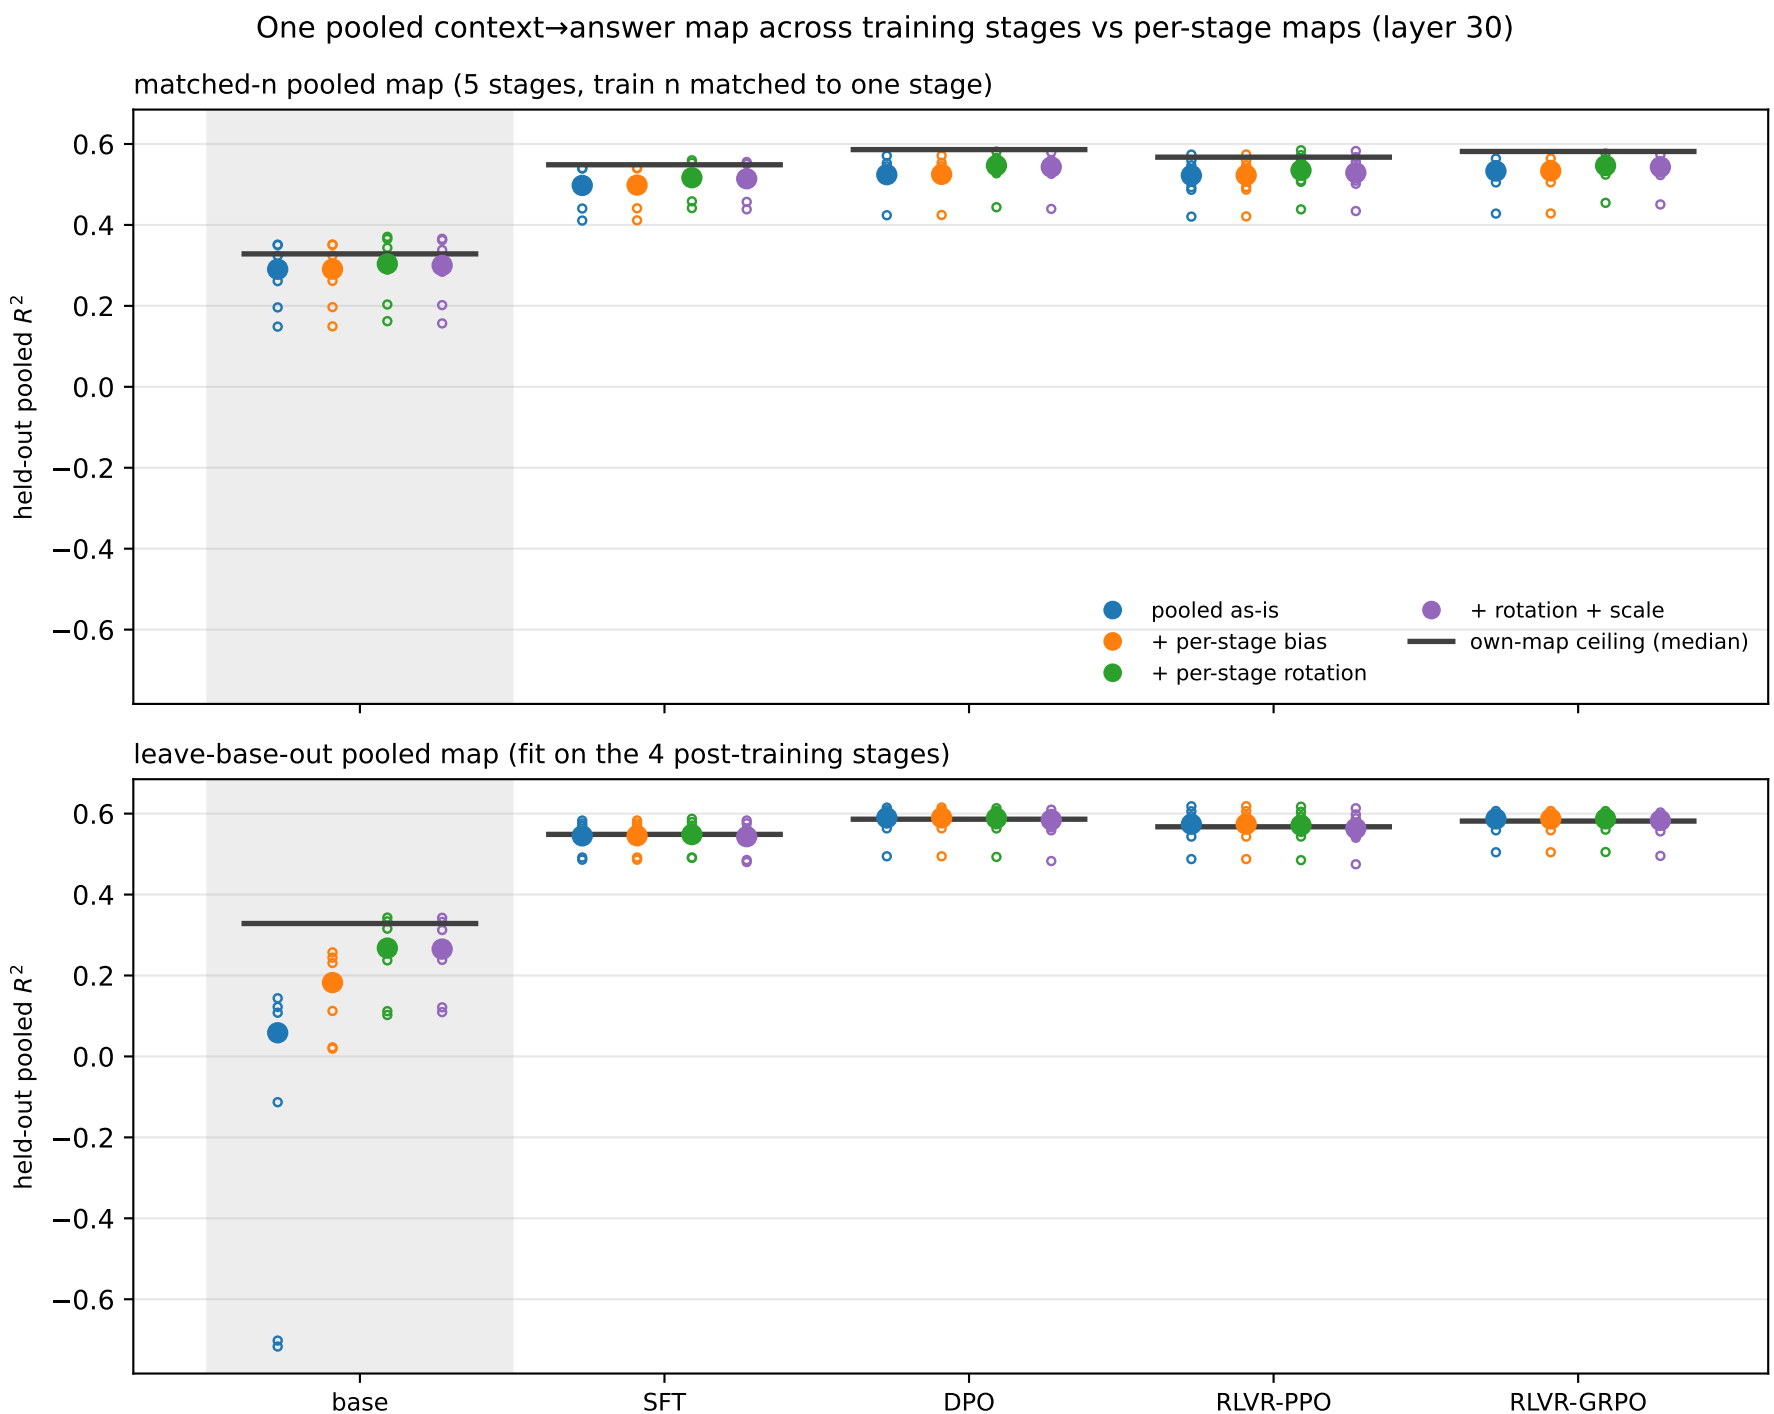
One pooled context→answer map across training stages vs per-stage maps (layer 30)
matched-n pooled map (5 stages, train n matched to one stage)
held-out pooled R^2
0.6
0.4
0.2
0.0
−0.2
−0.4
−0.6
pooled as-is
+ per-stage bias
+ per-stage rotation
+ rotation + scale
own-map ceiling (median)
leave-base-out pooled map (fit on the 4 post-training stages)
held-out pooled R^2
base
SFT
DPO
RLVR-PPO
RLVR-GRPO