

Which kinds of feature lose most when the map sees only the prefix, or only the query?
 $\text{drop} = R^2(\text{full context}) - R^2(\text{arm})$, per feature $\cdot n = 113,288 \cdot$ top 22 predictors by $|c - 0.5|$

