

Which kinds of feature lose most when the map sees only the prefix, or only the query?

drop = $R^2(\text{full context}) - R^2(\text{arm})$, per feature · n = 113,288 · context better on 65% of features vs prefix, 62% vs query-only · top 22 predictors by $|c - 0.5|$

