


The same predictors, controlling for average activity across tokens
pairs formed only within the 10 deciles of per-token firing frequency · sorted by $|c - 0.5|$ · n = 120,716
average activity across tokens itself scored 0.766 pooled



c = P(the feature with the higher predictor value is the better-predicted one), over pairs the predictor separates;
for a binary predictor this is the ordinary AUROC. The DIAMOND row (matryoshka tier) is scored on a DIFFERENT dictionary
(layer-20 SAE lens matryoshka, 16,384-feature panel), so it supports a same-direction claim but is NOT row-to-row comparable
on magnitude. Every value is MARGINAL — the other predictors are not controlled (pairwise correlations reach 0.85). No confidence intervals.
Rows are sorted STRONGEST first by $|c - 0.5|$; a lollipop pointing LEFT of 0.5 is inverted — that property marks features the map predicts WORSE.