


The same predictors, controlling for average activity across tokens  
pairs formed only within the 10 deciles of per-token firing frequency · sorted by  $|c - 0.5| \cdot n = 120,716$   
average activity across tokens itself scored 0.766 pooled



$c = P(\text{the feature with the higher predictor value is the better-predicted one})$ , over pairs the predictor separates;  
for a binary predictor this is the ordinary AUROC. The DIAMOND row (matryoshka tier) is scored on a DIFFERENT dictionary  
(layer-20 SAE lens matryoshka, 16,384-feature panel), so it supports a same-direction claim but is NOT row-to-row comparable  
on magnitude. Every value is MARGINAL — the other predictors are not controlled (pairwise correlations reach 0.85). No confidence intervals.  
Rows are sorted STRONGEST first by  $|c - 0.5|$ ; a lollipop pointing LEFT of 0.5 is inverted — that property marks features the map predicts WORSE.