


Round 0 — no controls — every pair counts
n = 128,450 features · winner: Variance explained in answer space (0.814)



Forward stepwise: each round adds the previous round's winner to the control set and re-scores the rest.
c = P(the feature with the higher predictor value is the better-predicted one), over pairs the predictor separates; for a binary predictor this is the ordinary AUROC. A lollipop LEFT of 0.5 is inverted. Bins per control variable are COARSENEDED as the control set grows, so later rounds condition less finely per variable and residual confounding grows with round number.
No confidence intervals — differences under ~0.02 should not be read.