

Prefix end state → answer: prediction quality by judged context category

TRUNCATION-CONTROLLED — the 978 responses that hit the 1024-token generation cap are EXCLUDED (n=8,947 of 9,925, 9.9%)

second number on each bar = shift from the all-rows figure · sorted best-predicted first · 95% paired bootstrap CI

Mean per-context R^2 (1 - normalized error)

