


Is per-context prediction quality just reading LENGTH? #1738 multi-turn holdout (n=9,941), L19 ridge



Equal-count deciles of each length variable; each point is the mean per-context R^2 of that decile. The prediction TARGET is a mean hidden state over the answer's tokens, so the answer-length panel is the mechanism's own variable: a longer answer averages over more tokens, shrinking target variance toward a genre-typical centre and becoming mechanically easier to hit. Titles give the POOLED Spearman; the within-category table in the sidebar is the composition-free read. No confidence intervals.