


What survives adjusting for firing frequency — and does the ordering differ between the linear and nonlinear maps?

dense context state (3,584-d) -> SAE answer features, 16,089 panel features (BatchTopK k=64, layer 19) | ridge median R^2 +0.177 (99% positive) vs MLP -0.028 (46%) | target and covariates BOTH #1482 single-turn — no cross-corpus join



rows sorted by $|\text{ridge partial } \rho|$; percentile 95% CI over 2,000 bootstrap draws (ranks fixed at the full sample) — at $n = 16,089$ they are often narrower than the markers.