


Effect modification: the same predictor behaves differently at different activity levels

113,262 SAE features, full dictionary (BatchTopK k=64, resid_post layer 19) | REAL TARGET — dense context -> SAE answer features, MLP(w=8192) mean pooling, full width. Same corpus as the covariates. The MLP has a heavy negative tail (median R² -0.877). | CI bands on the 4 most heterogeneous predictors only



FLAT => the single activity-partialled rho is a valid summary. SLOPED or SIGN-FLIPPING => effect modification; the profile replaces the single number. The RIGHTMOST point answers 'among only high-activating features, what explains predictability' — that decile is also the estimable slice (76.9% of its features have $R^2 > 0$), so measurement noise is absent there by construction.