



















Unified AUROC-at-depth: top-k vs bottom-k by R^2 , scored by each predictor

113,262 features | REAL TARGET — dense context -> SAE feature map, MLP(w=8192)___mean arm (#1482 single-turn corpus, 121,111 scored features; covariates same corpus — no cross-corpus mismatch)



grey band = per-k activity-stratified permutation null (2.5..97.5 pct). Blue = continuous predictor, green = binary one-vs-rest label as the score. Metric is AUROC-at-depth, never classification accuracy.