


Per-context predictions behind the sycophancy extremes (fair-allocation refit)

model-written evaluations (n=576 contexts, seed 0)



mimicry (n=285 contexts, seed 0)

